

Adam Wright

wrighada

CS-475

Project 4


1. What machine did I run this on?

I ran the simulation successfully on OSU's Flip server with one five and fifteen minute load averages of 3.66, 3.80, and 3.73 respectively. I originally attempted to run the tests on my local Windows machine, but I didn't get any resultant speedup from the tests.

2. Show the table of performances for each array size and the corresponding speedups.

| Array Size | Conventional Array MegaFlops/sec | SIMD MegaFlops/sec | SIMD Speedup |
|------------|-------------------------------------|-----------------------|--------------|
| 1000 | 818.40 | 3,508.96 | 4.29 |
| 8000 | 821.61 | 3,518.30 | 4.28 |
| 10,000 | 822.32 | 3,523.93 | 4.29 |
| 80,000 | 773.84 | 2,709.12 | 3.50 |
| 100,000 | 809.64 | 2,921.31 | 3.61 |
| 800,000 | 793.13 | 2,958.83 | 3.73 |
| 1,000,000 | 704.14 | 1,946.43 | 2.76 |
| 4,000,000 | 713.28 | 1318.97 | 1.85 |
| 8,000,000 | 705.55 | 897.59 | 1.27 |

3. Graph of SIMD/non-SIMD speedup versus array size



4. What patterns are you seeing in the speedups?

As the array size grows it is very clear that the algorithm is able to process the values in the array more quickly than the cache or memory bus are able to provide values. Before the 10,000 element mark the cache is well equipped to provide elements to process and between 80,000 and 800,000 elements there is only a slight decline in performance to a speedup of around 3.5. At 1,000,000 elements the speedup is below three and as it approaches 8,000,000 it is close to having no remaining speedup at all.

5. Are they consistent across a variety of array sizes?

No, the speedups decline significantly after the array size grows to larger than 1,000,000 elements. While the array size is below 10,000 elements, the performance is consistent, but between 80,000 and 1,000,000 elements there is a noticeable decline. The 4,000,000 element array does still approach a doubling of the speed, but by 8,000,000 there is nearly no speedup left at all.

6. Why or why not, do you think?


I believe that the algorithm is outrunning the supply of elements in the cache to provide elements for the registers. Once the registers need to make a round trip to RAM, the memory access negates the benefits of SIMD. This algorithm would benefit greatly from using prefetch, if the algorithm's primary purpose were to process arrays of more than 10,000 elements. Prefetch would likely allow the algorithm to maintain it's 4x speedup through the entire range of array sizes in this test.

EXTRA-CREDIT:

NUMT: Number of threads --- SIZE: array size --- Arr_MFL: non-SIMD array MegaFlops

Simd_MFL: SIMD MegaFlops --- S_UP-1: single threaded speedup --- S_MT_MFL: SIMD Multicore MegaFlops --- S_UP_2: multicore and SIMD speedup

| ***** ** Non-SIMD vs SIMD vs Multi-core SIMD ** ***** | | | | | | |
|---|---------|---------|----------|--------|----------|--------|
| NUMT | SIZE | Arr_MFL | Simd_MFL | S_UP-1 | S_MT_MFL | S_UP_2 |
| 1 | 1000 | 781.47 | 3334.60 | 4.27 | 1161.43 | 1.49 |
| 1 | 8000 | 821.77 | 3542.97 | 4.31 | 2769.16 | 3.37 |
| 1 | 10000 | 821.97 | 3547.80 | 4.32 | 2893.01 | 3.52 |
| 1 | 80000 | 809.24 | 2717.56 | 3.36 | 2606.80 | 3.22 |
| 1 | 100000 | 773.70 | 2780.99 | 3.59 | 2680.17 | 3.46 |
| 1 | 800000 | 729.17 | 2075.40 | 2.85 | 2039.45 | 2.80 |
| 1 | 1000000 | 779.08 | 2520.75 | 3.24 | 2647.73 | 3.40 |
| 1 | 4000000 | 713.28 | 1318.97 | 1.85 | 1301.80 | 1.83 |
| 1 | 8000000 | 706.77 | 1285.61 | 1.82 | 1302.12 | 1.84 |
| 2 | 1000 | 780.90 | 3334.60 | 4.27 | 756.96 | 0.97 |
| 2 | 8000 | 786.98 | 3373.89 | 4.29 | 2921.75 | 3.71 |
| 2 | 10000 | 787.03 | 3394.69 | 4.31 | 3115.46 | 3.96 |
| 2 | 80000 | 809.13 | 2805.84 | 3.47 | 5204.92 | 6.43 |
| 2 | 100000 | 809.28 | 2839.12 | 3.51 | 5007.66 | 6.19 |
| 2 | 800000 | 787.17 | 2964.86 | 3.77 | 5478.33 | 6.96 |
| 2 | 1000000 | 743.14 | 2689.74 | 3.62 | 4844.05 | 6.52 |
| 2 | 4000000 | 697.17 | 1292.38 | 1.85 | 1902.96 | 2.73 |
| 2 | 8000000 | 656.93 | 971.06 | 1.48 | 1198.82 | 1.82 |
| 3 | 1000 | 818.40 | 3508.96 | 4.29 | 589.97 | 0.72 |
| 3 | 8000 | 821.73 | 3514.70 | 4.28 | 3570.21 | 4.34 |
| 3 | 10000 | 786.97 | 3386.66 | 4.30 | 3831.37 | 4.87 |
| 3 | 80000 | 773.61 | 2749.44 | 3.55 | 7395.23 | 9.56 |
| 3 | 100000 | 773.81 | 2734.78 | 3.53 | 7303.37 | 9.44 |
| 3 | 800000 | 453.91 | 1713.87 | 3.78 | 4690.06 | 10.33 |
| 3 | 1000000 | 415.29 | 1515.80 | 3.65 | 4243.18 | 10.22 |
| 3 | 4000000 | 697.27 | 1236.39 | 1.77 | 2128.35 | 3.05 |
| 3 | 8000000 | 698.38 | 1244.02 | 1.78 | 1704.46 | 2.44 |
| 4 | 1000 | 819.03 | 3508.96 | 4.28 | 571.14 | 0.70 |
| 4 | 8000 | 786.73 | 3365.30 | 4.28 | 3420.24 | 4.35 |
| 4 | 10000 | 822.57 | 3547.21 | 4.31 | 4221.51 | 5.13 |
| 4 | 80000 | 773.21 | 2683.81 | 3.47 | 10743.46 | 13.89 |
| 4 | 100000 | 424.56 | 1599.44 | 3.77 | 5974.86 | 14.07 |
| 4 | 800000 | 734.29 | 2404.44 | 3.27 | 6760.08 | 9.21 |
| 4 | 1000000 | 449.18 | 2767.75 | 6.16 | 7650.00 | 17.03 |
| 4 | 4000000 | 701.73 | 1276.63 | 1.82 | 1929.37 | 2.75 |
| 4 | 8000000 | 640.70 | 1305.79 | 2.04 | 2067.03 | 3.23 |



E.C. A brief discussion of what the curves are showing and why I think it is working that way.

The curves are showing that in the regions where the cache can provide a sufficient number of elements to keep the registers filled, it is possible to get close to a 16 times improvement by doing SSE SIMD with four core multi-threading. The multicore version of the implementation still suffers from the same behavior as the single core SIMD implementation, which is that above 1,000,000 array elements, the speedup becomes much smaller. The reason for this is the same as the reason for the performance degradation in the single core version which is that the cache is unable to provide contain enough values for the algorithm to remain efficient. The addition of properly implemented prefetch would allow the larger array sizes to remain performant, though at array sizes below 1,000,000 this non-prefetch implementation is a great performer.